

DTGen ASOF Demonstration

Developed by DMSTEX (<http://dmstex.com>)

Oracle and Java are registered trademarks of Oracle Corporation and/or its affiliates.

Table of Contents

Introduction.....	1
Exercise #1: Entity Based History and Audit.....	1
Exercise #2: EFF vs. LOG Table Types.....	5
Exercise #3: Point-in-Time ASOF Views.....	10
Exercise #4: Audited POP Functions.....	10
Exercise #5: Comprehensive OMNI Views.....	10
Exercise #6: Transportable ASOF Data.....	11

Introduction:

The exercises in this demonstration are focused on the history and audit functionality of DTGen functionality. All functionality in these exercises is available through both command line and graphical user interface (GUI) mode. For simplicity in understanding the under-lying workings of DTGen, these exercises are conducted entirely in command-line mode.

The "basic" demonstration should be reviewed before running these exercises. Several concepts introduced in those exercises not explained here. Exercise #1 in this demonstration is similar to Exercise #1 in the basic demonstration.

The exercises in this directory are numbered and must be executed in sequential order. The demonstration users must be created with the "create_demo_users.sql" script in the parent directory before the first exercise is run. The demonstration users must be dropped with the "drop_demo_users.sql" script before the "create_demo_users.sql" script can be re-run. These exercises also assume that the default username/password (dtgen/dtgen) is still in use for the generator. Names and passwords are set in the "vars.sql" script and can be modified, if necessary. Also, the DTGen database objects must be installed in the database and the DTGen must be ready to generate code.

Exercise #1: Entity Based History and Audit

Command Line:

```
sqlplus /nolog @e1
```

Exercise #1 modifies the database. The "drop_demo_users.sql" and "create_demo_users.sql" scripts must be used to reset the database before re-running this exercise.

Based on the demobld.sql script, this exercise implements the EMP and DEPT tables using DTGen. The script for this exercise performs the following functions:

1. Removes any old DEMO2 Items from DTGEN
2. Creates new DEMO2 Items in DTGEN
3. Generates the DEMO2 Application in DTGEN
4. Creates the "install_db.sql" script
5. Runs the "install_db.sql" script
6. Loads and Reports Data

Steps 1-3 are captured in the "e1.LST" file. Following is a example of e1.LST.

```
Login to dtgen
Connected.
Remove old DEMO2 Schema from DTGEN
create a DEMO2 Schema in DTGEN
Generate Demo2 Application
Capture install_db.sql Script
```

Step 4 is captured in the "install_db.sql" file. This file is about 79 kbytes and has over 3,000 lines. Due to its size, it is not listed here. It contains all the code generated by DTGen for this application.

Steps 5 and 6 are captured in the "install.LST" file. Step 5 is the execution of the install_db.sql script.

```
Login to dtgen_db_demo
Connected.
```

```
FILE_NAME
-----
-) create_glob
```

```
FILE_NAME
-----
-) create_ods
```

```
TABLE_NAME
-----
*** dept ***
```

```
TABLE_NAME
-----
*** emp ***
```

```
FILE_NAME
-----
-) create_integ
```

```
TABLE_NAME
-----
*** dept ***
```

```
TABLE_NAME
-----
*** emp ***
```

```
FILE_NAME
-----
-) create_oltp
```

```
TABLE_NAME
-----
*** dept ***
```

```

TABLE_NAME
-----
***   emp   ***

FILE_NAME
-----
-) create_mods

```

The above listing represents a successful installation of the application generated by DTGen. This application is small in that it only has 2 tables, 1 tier (the database tier), and no user schema.

The DEPT table is silently loaded with data. A query of column comments on the DEPT table from the data dictionary help identify what each column's data represents. Following the column comments is a report of all the data in the DEPT table (active view) for the selected columns.

COLUMN_NAME	COMMENTS
DEPTNO	Department Number
DNAME	Name of the Department
LOC	Location for the Department
AUD_BEG_USR	User that created this record
AUD_BEG_DTM	Date/Time this record was created (must be in nanoseconds)

DEPTNO	DNAME	LOC	AUD_BEG_USR	AUD_BEG_DTM
10	ACCOUNTING	NEW YORK	Dataload	01-NOV-80 12
20	RESEARCH	DALLAS	Dataload	01-NOV-80 12
30	SALES	CHICAGO	THOMPSON	17-AUG-81 12
40	OPERATIONS	BOSTON	JAMES	12-FEB-82 12

The EMP table is also silently loaded with data. The same queries of column comments and data on the EMP table (active view) are shown.

COLUMN_NAME	COMMENTS
EMPNO	Employee Number
ENAME	Employee Name
JOB	Job Title
MGR_EMP_NK1	EMP Natural Key Value 1: Employee Number
HIREDATE	Date the Employee was hired
SAL	Employee's Salary
DEPT_NK1	DEPT Natural Key Value 1: Department Number
EFF_BEG_DTM	Date/Time this record became effective
AUD_BEG_USR	User that created this record

EMPNO	ENAME	JOB	MGR_	HIREDATE	SAL	DEPT_	AUD_BEG_USR	AUD_BEG_DTM
7369	SMITH	CLERK	7902	17-DEC-80	800	20	SMITH	27-FEB-83 12
7499	ALLEN	SALESMAN	7698	20-FEB-81	1600	30	THOMPSON	15-MAY-81 12
7521	WARD	SALESMAN	7698	22-FEB-81	1250	30	THOMPSON	16-MAY-81 12
7566	JONES	MANAGER	7839	02-APR-81	2975	20	SMITH	01-DEC-81 12
7654	MARTIN	SALESMAN	7698	28-SEP-81	1250	30	SMITH	28-SEP-81 12
7698	BLAKE	MANAGER	7839	01-MAY-81	2850	30	SMITH	27-NOV-81 12
7782	CLARK	MANAGER	7839	09-JUN-81	2450	10	SMITH	26-NOV-81 12
7788	SCOTT	ANALYST	7566	09-DEC-82	3000	20	JAMES	07-DEC-82 12
7839	KING	PRESIDENT		17-NOV-81	5000	10	SMITH	19-NOV-81 12
7844	TURNER	SALESMAN	7698	08-SEP-81	1500	30	SMITH	09-SEP-81 12
7876	ADAMS	CLERK	7788	12-JAN-83	1100	20	SMITH	12-JAN-83 12
7900	JAMES	CLERK	7698	03-DEC-81	950	30	SMITH	04-DEC-81 12
7902	FORD	ANALYST	7566	03-DEC-81	3000	20	SMITH	05-DEC-81 12
7934	MILLER	CLERK	7782	23-JAN-82	1300	10	JAMES	23-JAN-82 12

In addition to the DEPT and EMP tables, this exercise also loaded a DEPT audit table called DEPT_AUD and an EMP history table called EMP_HIST.

COLUMN_NAME	COMMENTS
DEPT_ID	Surrogate Primary Key from the ACTIVE table
DEPTNO	Department Number
DNAME	Name of the Department
LOC	Location for the Department
AUD_BEG_USR	User that created this record
AUD_BEG_DTM	Date/Time this record was created (must be in nanoseconds)
AUD_END_USR	User that modified/deleted this record
AUD_END_DTM	Date/Time this record was modified/deleted (must be in nanoseconds)

DEPT_ID	DEPTNO	DNAME	LOC	AUD_BEG_	AUD_BEG_D	AUD_END_	AUD_END_D
3	20	SALES	ST LOUIS	Datload	01-NOV-80	THOMPSON	17-AUG-82
4	40	OPERATIONS	BUFFALO	Datload	01-NOV-80	JAMES	12-FEB-82

In the DEPT_AUD listing above, the column comments for the columns starting with "AUD_" indicate the user and date/time of modifications are being tracked, even after a record has been updated or deleted. The data records show that the SALES department was originally in St. Louis and the OERATIONS department was originally in Buffalo.

COLUMN_NAME	COMMENTS
EMP_ID	Surrogate Primary Key from the ACTIVE table
EMPNO	Employee Number
ENAME	Employee Name
JOB	Job Title
MGR_EMP_ID	Surrogate Key of Employee's Manager
AUD_BEG_USR	User that created this record
AUD_BEG_DTM	Date/Time this record was created (must be in nanoseconds)
AUD_END_USR	User that modified/deleted this record
AUD_END_DTM	Date/Time this record was modified/deleted (must be in nanoseconds)

EMP_ID	EMPNO	ENAME	JOB	MGR_	AUD_BEG_	AUD_BEG_D	AUD_END_	AUD_END_D
1	7301	ELLISON	PRESIDENT		DAVIS	31-OCT-80	THOMPSON	28-JUN-81
2	7344	DAVIS	CLERK	1	DAVIS	15-NOV-80	THOMPSON	21-JUN-81
2	7344	DAVIS	CLERK	11	THOMPSON	22-JUN-81	THOMPSON	20-AUG-81
2	7344	DAVIS	CLERK	12	SMITH	21-AUG-81	SMITH	30-NOV-81
2	7344	DAVIS	CLERK	15	SMITH	28-NOV-81	SMITH	10-DEC-81
3	7369	THOMPSON	CLERK	1	DAVIS	18-DEC-80	SMITH	25-FEB-83
3	7369	SMITH	CLERK	15	SMITH	29-NOV-81	SMITH	28-FEB-83
3	7369	THOMPSON	CLERK	11	THOMPSON	22-JUN-81	SMITH	22-AUG-81
3	7369	SMITH	CLERK	12	SMITH	23-AUG-81	SMITH	30-NOV-81
4	7499	ALLEN	SALESMAN	1	THOMPSON	18-FEB-81	THOMPSON	12-MAY-81
5	7521	WARD	SALESMAN	1	THOMPSON	24-FEB-81	THOMPSON	16-MAY-81
6	7566	JONES	MANAGER	1	THOMPSON	31-MAR-81	THOMPSON	21-JUN-81
6	7566	JONES	MANAGER	11	THOMPSON	23-JUN-81	THOMPSON	19-AUG-81
6	7566	JONES	MANAGER	12	THOMPSON	20-AUG-81	SMITH	29-NOV-81
7	7654	MARTIN	SALESMAN	1	THOMPSON	17-APR-81	THOMPSON	14-MAY-81
8	7698	BLAKE	MANAGER	1	THOMPSON	30-APR-81	THOMPSON	24-JUN-81
8	7698	BLAKE	MANAGER	11	THOMPSON	21-JUN-81	SMITH	23-AUG-81
8	7698	BLAKE	MANAGER	12	SMITH	22-AUG-81	SMITH	29-NOV-81
9	7782	CLARK	MANAGER	1	THOMPSON	09-JUN-81	THOMPSON	24-JUN-81
9	7782	CLARK	MANAGER	11	THOMPSON	22-JUN-81	SMITH	22-AUG-81
9	7782	CLARK	MANAGER	12	SMITH	21-AUG-81	SMITH	27-NOV-81
10	7788	SCOTT	ANALYST	6	THOMPSON	09-JUN-81	JAMES	09-MAR-82
11	7839	KING	PRESIDENT		THOMPSON	18-JUN-81	SMITH	30-AUG-81
12	7840	LANE	PRESIDENT		THOMPSON	12-AUG-81	SMITH	28-NOV-81
16	7876	ADAMS	CLERK	6	SMITH	22-NOV-81	JAMES	14-JUN-82

In the EMP_HIST listing above, the column comments and data for selected columns are queried. The column comments for the columns show that both "_AUD" and "_HIST" tables include the audit columns that start with "AUD_". The data records show that many changes have occurred to the data in the EMP table.

Exercise #2: EFF vs. LOG Table Types

Command Line:

```
sqlplus /nolog @e2
```

Exercise #2 does not modify the database. This exercise can be repeated without problem.

Exercise #1 included a brief introduction to the DEPT_AUD and EMP_HIST tables. Below are the DTGen settings that generated those tables.

```
Login to dtgen
Connected.
```

```
VALUE DESCRIPTION
-----
EFF  A historical table type with effective/audit begin/end timestamps and
      begin/end user recording
LOG  An audit table type with audit only begin/end timestamps and begin/end
      user recording
NON  A none or nothing table type without begin/end timestamps or begin/end
      user recording
```



```
      SEQ NAME          TYP
-----
      10 dept           LOG
      20 emp            EFF
```

For any table configured in DTGen, one of these 3 table types must be selected. In the "basics" demonstration, "NON" was set as the table type for both tables. This demonstration shows the use of LOG and EFF table types. An example of when to use LOG versus EFF is also represented with these tables. The DEPT table holds department information, which is slow moving (doesn't change often or rapidly) and is generally known well in advance of any changes to the applicaiton data (i.e. Adding a new department). Constrast the employee information, which can be fast moving and needs to be recorded as occuring at a specific time (effectivity) in addition to simple audit recording.

```
SQL>
SQL> select id, deptno dept, loc,
2      aud_beg_usr, aud_beg_dtm
3      from dept_act where deptno = 50;

no rows selected

SQL> select dept_id id, deptno dept, loc,
2      aud_beg_usr, aud_beg_dtm, aud_end_usr, aud_end_dtm
3      from dept_aud where deptno = 50;

no rows selected
```

There are no records for a department with a DEPTNO of "50".

```
SQL>
SQL> execute util.set_usr('USER1');

PL/SQL procedure successfully completed.

SQL> select systimestamp from dual;

SYSTIMESTAMP
-----
17-APR-12 04.35.54
```

```

1 row selected.

SQL> insert into dept_act (deptno, dname, loc)
2   values (50, 'NEW_DEPT', 'LZ');

1 row created.

SQL>
SQL> select id, deptno dept, loc,
2   aud_beg_usr, aud_beg_dtm
3   from dept_act where deptno = 50;

ID DEPT LOC AUD_BEG_USR AUD_BEG_DTM
---
5   50 LZ   USER1       17-APR-12 04.35.54

1 row selected.

SQL> select dept_id id, deptno dept, loc,
2   aud_beg_usr, aud_beg_dtm, aud_end_usr, aud_end_dtm
3   from dept_aud where deptno = 50;

no rows selected

```

After setting the util.set_usr, the time is reported and a DEPT record is created. The last 2 queries above show that a DEPT_ACT record was created by the util.set_usr value at the current time, and here are no DEPT_AUD records.

```

SQL>
SQL> execute dbms_lock.sleep(1);

PL/SQL procedure successfully completed.

SQL> execute util.set_usr('USER2');

PL/SQL procedure successfully completed.

SQL> select systimestamp from dual;

SYSTIMESTAMP
-----
17-APR-12 04.35.55

1 row selected.

SQL> update dept_act
2   set loc = 'LA'
3   where deptno = 50;

1 row updated.

SQL>
SQL> select id, deptno dept, loc,
2   aud_beg_usr, aud_beg_dtm
3   from dept_act where deptno = 50;

ID DEPT LOC AUD_BEG_USR AUD_BEG_DTM
---
5   50 LA   USER2       17-APR-12 04.35.55

1 row selected.

SQL> select dept_id id, deptno dept, loc,
2   aud_beg_usr, aud_beg_dtm, aud_end_usr, aud_end_dtm
3   from dept_aud where deptno = 50;

ID DEPT LOC AUD_BEG_USR AUD_BEG_DTM AUD_END_USR AUD_END_DTM
---
5   50 LZ   USER1       17-APR-12 04.35.54 USER2       17-APR-12 04.35.55

1 row selected.

```

In the sequence above, the DBMS_LOCK.SLEEP function is used to elapse one second before

running the update with a different util.set_usr. The last 2 queries above show one record in each DEPT_ACT and DEPT_AUD. The DEPT_AUD record is the previous DEPT_ACT record. Also, the DEPT_AUD record has a matching AUD_END_DTM to the AUD_BEG_DTM of the new DEPT_ACT record. Notice that USER2 is recorded as the AUD_END_USR in DEPT_AUD and the AUD_BEG_USR in DEPT_ACT.

```
SQL>
SQL> execute dbms_lock.sleep(1);

PL/SQL procedure successfully completed.

SQL> execute util.set_usr('USER3');

PL/SQL procedure successfully completed.

SQL> select systimestamp from dual;

SYSTIMESTAMP
-----
17-APR-12 04.35.56

1 row selected.

SQL> delete from dept_act where deptno = 50;

1 row deleted.

SQL>
SQL> select id, deptno dept, loc,
2     aud_beg_usr, aud_beg_dtm
3     from dept_act where deptno = 50;

no rows selected

SQL> select dept_id id, deptno dept, loc,
2     aud_beg_usr, aud_beg_dtm, aud_end_usr, aud_end_dtm
3     from dept_audit where deptno = 50;
```

ID	DEPT	LOC	AUD_BEG_USR	AUD_BEG_DTM	AUD_END_USR	AUD_END_DTM
5	50	LZ	USER1	17-APR-12 04.35.54	USER2	17-APR-12 04.35.55
5	50	LA	USER2	17-APR-12 04.35.55	USER3	17-APR-12 04.35.56

```

2 rows selected.

```

In the sequence above, a delete is run on DEPT_ACT, resulting in no DEPT_ACT records and 2 DEPT_AUD records. This functionality works the same for the DEPT_DML API calls.

Below, the same set of processes will be repeated for an employee record.

```
SQL>
SQL> select id, empno, ename, to_char(eff_beg_dtm,'DD HH24:MI:SS') eff_beg_dtm,
2     aud_beg_usr, to_char(aud_beg_dtm,'DD HH24:MI:SS') aud_beg_dtm
3     from emp_act where empno = 9999;

no rows selected

SQL> select emp_id id, empno, ename,
2     to_char(eff_beg_dtm,'DD HH24:MI:SS') eff_beg_dtm,
3     aud_beg_usr, to_char(aud_beg_dtm,'DD HH24:MI:SS') aud_beg_dtm,
4     to_char(eff_end_dtm,'DD HH24:MI:SS') eff_end_dtm,
5     aud_end_usr, to_char(aud_end_dtm,'DD HH24:MI:SS') aud_end_dtm
6     from emp_hist where empno = 9999;

no rows selected

SQL>
SQL> execute util.set_usr('USER1');

PL/SQL procedure successfully completed.
```

```

SQL> select systimestamp from dual;

SYSTIMESTAMP
-----
17-APR-12 08.25.17

1 row selected.

SQL> insert into emp_act (empno, ename, job, hiredate, sal, dept_nk1,
2     eff_beg_dtm)
3     values (9999, 'NEW_EMP', 'CLERK', sysdate, 100, 40,
4     to_timestamp('1983-6-1 11', 'YYYY-MM-DD HH24'));

1 row created.

SQL>
SQL> select id, empno, ename, to_char(eff_beg_dtm,'DD HH24:MI:SS') eff_beg_dtm,
2     aud_beg_usr, to_char(aud_beg_dtm,'DD HH24:MI:SS') aud_beg_dtm
3     from emp_act where empno = 9999;

ID EMPNO ENAME      EFF_BEG_DTM AUD_B AUD_BEG_DTM
-----
22  9999 NEW_EMP    01 11:00:00 USER1 17 20:25:17

1 row selected.

SQL> select emp_id id, empno, ename,
2     to_char(eff_beg_dtm,'DD HH24:MI:SS') eff_beg_dtm,
3     aud_beg_usr, to_char(aud_beg_dtm,'DD HH24:MI:SS') aud_beg_dtm,
4     to_char(eff_end_dtm,'DD HH24:MI:SS') eff_end_dtm,
5     aud_end_usr, to_char(aud_end_dtm,'DD HH24:MI:SS') aud_end_dtm
6     from emp_hist where empno = 9999;

no rows selected

```

Notice that the EMP_ACT active view has the EFF_BEG_DTM column that DEPT_ACT does not. This is a key difference between EFF and LOG. This additional effectivity column allows a date/time in the past to be entered and referenced as the effective time an event occurred. Otherwise, the LOG table only records when an event was entered, as the AUD_BEG_DTM is acquired from the system clock and cannot be entered.

The first 2 queries show there are no EMP_ACT records with an EMPNO of 9999. After the user is set and the time is queried, a record is inserted into the EMP_ACT active view with an effectivity date/time of June 1st, 1983 at 11am. The record is returned from the EMP_ACT query and no records are returned from EMP_HIST.

```

SQL>
SQL> execute dbms_lock.sleep(1);

PL/SQL procedure successfully completed.

SQL> execute util.set_usr('USER2');

PL/SQL procedure successfully completed.

SQL> select systimestamp from dual;

SYSTIMESTAMP
-----
17-APR-12 08.25.18

1 row selected.

SQL> update emp_act
2     set ename = 'UPD_EMP',
3     eff_beg_dtm = to_timestamp('1983-6-2 12', 'YYYY-MM-DD HH24')
4     where empno = 9999;

1 row updated.

SQL>

```



```
SQL> select id, empno, ename, to_char(eff_beg_dtm,'DD HH24:MI:SS') eff_beg_dtm,
2   aud_beg_usr, to_char(aud_beg_dtm,'DD HH24:MI:SS') aud_beg_dtm
3   from emp_act where empno = 9999;
```

ID	EMPNO	ENAME	EFF_BEG_DTM	AUD_B	AUD_BEG_DTM
22	9999	UPD_EMP	02 12:00:00	USER2	17 20:25:18

1 row selected.

```
SQL> select emp_id id, empno, ename,
2   to_char(eff_beg_dtm,'DD HH24:MI:SS') eff_beg_dtm,
3   aud_beg_usr, to_char(aud_beg_dtm,'DD HH24:MI:SS') aud_beg_dtm,
4   to_char(eff_end_dtm,'DD HH24:MI:SS') eff_end_dtm,
5   aud_end_usr, to_char(aud_end_dtm,'DD HH24:MI:SS') aud_end_dtm
6   from emp_hist where empno = 9999;
```

ID	EMPNO	ENAME	EFF_BEG_DTM	AUD_B	AUD_BEG_DTM	EFF_END_DTM	AUD_E	AUD_END_DTM
22	9999	NEW_EMP	01 11:00:00	USER1	17 20:25:17	02 12:00:00	USER2	17 20:25:18

1 row selected.

Again, the user is set and the time is queried. An update is issued against the active view with an effectivity date/time of June 2nd, 1983 at 12pm. Each of the last 2 queries return 1 record showing the effectivity that was entered. AUD_BEG_DTM and AUD_END_DTM will always have the current date/time as these values are not allowed to be entered.

```
SQL>
SQL> execute dbms_lock.sleep(1);

PL/SQL procedure successfully completed.
```

```
SQL> select systimestamp from dual;
```

SYSTIMESTAMP
17-APR-12 08.25.19

1 row selected.

```
SQL> declare
2   eff_end_dtm timestamp with local time zone;
3   emp_id      number;
4   begin
5   util.set_usr('USER3');
6   select id into emp_id
7   from emp_act where empno = 9999;
8   eff_end_dtm := to_timestamp('1983-6-3 13', 'YYYY-MM-DD HH24');
9   emp_dml.del(emp_id, eff_end_dtm);
10  end;
11  /
```

PL/SQL procedure successfully completed.

```
SQL>
SQL> select id, empno, ename, to_char(eff_beg_dtm,'DD HH24:MI:SS') eff_beg_dtm,
2   aud_beg_usr, to_char(aud_beg_dtm,'DD HH24:MI:SS') aud_beg_dtm
3   from emp_act where empno = 9999;
```

no rows selected

```
SQL> select emp_id id, empno, ename,
2   to_char(eff_beg_dtm,'DD HH24:MI:SS') eff_beg_dtm,
3   aud_beg_usr, to_char(aud_beg_dtm,'DD HH24:MI:SS') aud_beg_dtm,
4   to_char(eff_end_dtm,'DD HH24:MI:SS') eff_end_dtm,
5   aud_end_usr, to_char(aud_end_dtm,'DD HH24:MI:SS') aud_end_dtm
6   from emp_hist where empno = 9999;
```

ID	EMPNO	ENAME	EFF_BEG_DTM	AUD_B	AUD_BEG_DTM	EFF_END_DTM	AUD_E	AUD_END_DTM
22	9999	NEW_EMP	01 11:00:00	USER1	17 20:25:17	02 12:00:00	USER2	17 20:25:18
22	9999	UPD_EMP	02 12:00:00	USER2	17 20:25:18	03 13:00:00	USER3	17 20:25:19

```
2 rows selected.
```

```
SQL>
SQL> rollback;
```

```
Rollback complete.
```

Notice the EMP_DML.DEL procedure was used to delete the record using an effectivity date/time of June 3rd, 1983 at 1pm. The DML package API is the only way to enter an EFF_END_DTM during a delete. There is no corresponding functionality available using the EMP_ACT active view. Otherwise, the EMP_DML.INS and EMP_DML_UPD procedures work the same as insert and delete SQL on the EMP_ACT active view.

Exercise #3: Point-in-Time ASOF Views

Command Line:

```
sqlplus /nolog @e3
```

Exercise #3 does not modify the database. This exercise can be repeated without problem.

In this exercise, indexes on foreign keys and natural keys are explored. Following is a query of the DTGen setup used to generate this application

```
SQL>
SQL> execute util.set_asof_dtm(to_timestamp('1981-06-01', 'YYYY-MM-DD'))
```

```
PL/SQL procedure successfully completed.
```

```
SQL>
SQL> select id, empno, ename, job, mgr_emp_nk1, hiredate,
2         sal, dept_nk1, aud_beg_usr, aud_beg_dtm
3        from emp_asof
4        order by empno;
```

ID	EMPNO	ENAME	JOB	MGR_	HIREDATE	SAL	DEPT	AUD_BEG_USR	AUD_BEG_DTM
1	7301	ELLISON	PRESIDENT		02-NOV-80	4000	10	DAVIS	31-OCT-80 12
2	7344	DAVIS	CLERK	7301	16-NOV-80	1400	10	DAVIS	15-NOV-80 12
3	7369	THOMPSON	CLERK	7301	17-DEC-80	800	10	DAVIS	18-DEC-80 12
4	7499	ALLEN	SALESMAN	7698	20-FEB-81	1600	30	THOMPSON	15-MAY-81 12
5	7521	WARD	SALESMAN	7698	22-FEB-81	1250	30	THOMPSON	16-MAY-81 12
6	7566	JONES	MANAGER	7301	02-APR-81	2975	20	THOMPSON	31-MAR-81 12
8	7698	BLAKE	MANAGER	7301	01-MAY-81	2850	30	THOMPSON	30-APR-81 12

```
7 rows selected.
```

```
SQL>
SQL> execute util.set_asof_dtm(to_timestamp('1981-09-01', 'YYYY-MM-DD'))
```

```
PL/SQL procedure successfully completed.
```

```
SQL>
SQL> select id, empno, ename, job, mgr_emp_nk1, hiredate,
2         sal, dept_nk1, aud_beg_usr, aud_beg_dtm
3        from emp_asof
4        order by empno;
```

ID	EMPNO	ENAME	JOB	MGR_	HIREDATE	SAL	DEPT	AUD_BEG_USR	AUD_BEG_DTM
2	7344	DAVIS	CLERK	7840	16-NOV-80	1400	10	SMITH	21-AUG-81 12
3	7369	SMITH	CLERK	7840	17-DEC-80	800	10	SMITH	23-AUG-81 12
3	7369	THOMPSON	CLERK		17-DEC-80	800	10	DAVIS	18-DEC-80 12
4	7499	ALLEN	SALESMAN	7698	20-FEB-81	1600	30	THOMPSON	15-MAY-81 12
5	7521	WARD	SALESMAN	7698	22-FEB-81	1250	30	THOMPSON	16-MAY-81 12
6	7566	JONES	MANAGER	7840	02-APR-81	2975	20	THOMPSON	20-AUG-81 12

```

      8 7698 BLAKE      MANAGER      7840 01-MAY-81 2850 30 SMITH      22-AUG-81 12
      9 7782 CLARK      MANAGER      7840 09-JUN-81 2450 10 SMITH      21-AUG-81 12
     10 7788 SCOTT      ANALYST      7566 12-JUN-81 3000 20 THOMPSON    09-JUN-81 12
     12 7840 LANE      PRESIDENT      14-AUG-81 6000 10 THOMPSON    12-AUG-81 12

```

10 rows selected.

SQL>

SQL> execute util.set_asof_dtm(to_timestamp('1982-01-01', 'YYYY-MM-DD'))

PL/SQL procedure successfully completed.

SQL>

```

SQL> select id, empno, ename, job, mgr_emp_nk1, hiredate,
2      sal, dept_nk1, aud_beg_usr, aud_beg_dtm
3      from emp_asof
4      order by empno;

```

ID	EMPNO	ENAME	JOB	MGR_	HIREDATE	SAL	DEPT	AUD_BEG_USR	AUD_BEG_DTM
3	7369	SMITH	CLERK	7839	17-DEC-80	800	10	SMITH	29-NOV-81 12
3	7369	THOMPSON	CLERK		17-DEC-80	800	10	DAVIS	18-DEC-80 12
4	7499	ALLEN	SALESMAN	7698	20-FEB-81	1600	30	THOMPSON	15-MAY-81 12
5	7521	WARD	SALESMAN	7698	22-FEB-81	1250	30	THOMPSON	16-MAY-81 12
6	7566	JONES	MANAGER	7839	02-APR-81	2975	20	SMITH	01-DEC-81 12
14	7654	MARTIN	SALESMAN	7698	28-SEP-81	1250	30	SMITH	28-SEP-81 12
8	7698	BLAKE	MANAGER	7839	01-MAY-81	2850	30	SMITH	27-NOV-81 12
9	7782	CLARK	MANAGER	7839	09-JUN-81	2450	10	SMITH	26-NOV-81 12
10	7788	SCOTT	ANALYST	7566	12-JUN-81	3000	20	THOMPSON	09-JUN-81 12
15	7839	KING	PRESIDENT		17-NOV-81	5000	10	SMITH	19-NOV-81 12
13	7844	TURNER	SALESMAN	7698	08-SEP-81	1500	30	SMITH	09-SEP-81 12

ID	EMPNO	ENAME	JOB	MGR_	HIREDATE	SAL	DEPT	AUD_BEG_USR	AUD_BEG_DTM
16	7876	ADAMS	CLERK	7566	22-NOV-81	1100	20	SMITH	22-NOV-81 12
17	7900	JAMES	CLERK	7698	03-DEC-81	950	30	SMITH	04-DEC-81 12
18	7902	FORD	ANALYST	7566	03-DEC-81	3000	20	SMITH	05-DEC-81 12

14 rows selected.

Exercise #4: Audited POP Functions

Command Line:

```
sqlplus /nolog @e4
```

Exercise #4 modifies the database. The "drop_demo_users.sql", "create_demo_users.sql", and "e1.sql" scripts must be used to reset the database before re-running this exercise.

Each table defined in DTGen is generated with a corresponding "active view". The DEPT and EMP tables have an active view called "DEPT_ACT" and "EMP_ACT", respectively. In most cases, these views should be used for all DML (Data

Exercise #5: Comprehensive OMNI Views

Command Line:

```
sqlplus /nolog @e5
```

Exercise #5 does not modify the database. This exercise can be repeated without problem.

The EMP table has a self-referencing foreign key. It is the relationship between employees and managers. Since managers are also employees, they have managers as well, with the exception of

the PRESIDENT. This self-referencing foreign

Exercise #6: Transportable ASOF Data

Command Line:

```
sqlplus /nolog @e6
```

Exercise #6 does not modify the database. This exercise can be repeated without problem.

Unlike the original demobld.sql, this demonstration includes built in domain checking on the JOB column in the EMP table. The configuration of DTGen included a domain specification for all possible company jobs. Unlike a foreign key table, a domain is embedded into the error checking

Oracle introduced a mechanism to hold rollback entries for use by "flashback" queries to provide data as it was at some previous point in time. Since the source of this data is the rollback segments, transporting this data has limitations. DTGen provides a schema based approach to retrieving data at some previous point in time. The data in the Dtgen generated schema are more easily transported to different databases and servers than the rollback segment data. The loading of data for this demonstration demonstrates some of the capability of transporting historical data. These exercises would not have been more difficult and abstract using Oracle's rollback segment data.

Additionally, it is impossible to edit historical mistakes in the rollback segments.