

DTGen Tiers Demonstration

Developed by DMSTEX (<http://dmstex.com>)

Oracle and Java are registered trademarks of Oracle Corporation and/or its affiliates.

Table of Contents

Introduction.....	1
Exercise #1: Simple Mid-Tier.....	2
Exercise #2: Materialized Views.....	6
Exercise #3: User Security.....	8
Exercise #4: Global Locks.....	9

Introduction:

The exercises in this demonstration are focused on DTGen functionality that enables and enhances tiered deployment. All functionality in these exercises is available through both command line and graphical user interface (GUI) mode. For simplicity in understanding the under-lying workings of DTGen, these exercises are conducted entirely in command-line mode.

Tiered deployment as described here includes 2 forms of deployment:

- database/mid-tier server tiers
- database user/schema tiers

Multi-tiered hardware deployment is a common aspect of many systems. The database and mid-tier servers use a software deployment strategy that increases capacity and improves security. Capacity can be increased by adding servers to existing systems instead of replacing centralized servers with larger servers. Common deployment architectures used by Oracle for the Transaction Performance Processing Council's TPC-C Benchmark (<http://www.tpc.org/tpcc/default.asp>) were used as a model for DTGen's multi-tiered deployment in that all possible functionality is moved from the database tier to the mid-tier.

Security can be improved through the user of layered access to applications and data. Layered security can also be found in many systems. Oracle e-Business Suite layers application schema behind a common database application login. The layered security generated by DTGen was created to assist in compliance with "The Defense Information Services (DISA) Oracle Database Security Readiness Review (SRS)"

("http://iase.disa.mil/stigs/downloads/zip/unclassified_oracle10_v8r1.8_checklist_20100827.zip")
Specific references in the "U_DB_oracle10_v8r1.8_Checklist_20100827.pdf" file include:

- V0005683 - Application object owner accounts should be disabled when not performing

installation or maintenance actions.

- V0015613 - Each database user, application or process should have an individually assigned account.
- V0015629 - Application users privileges should be restricted to assignment using application user roles.
- V0003847 - Database application user accounts should be denied storage usage for object creation within the database.

The exercises in this demonstration do not entirely conform to the DISA Oracle Database SRS because they use only 1 database and there is only 1 configuration for the application. Full compliance is obtained with multiple databases in that the same shema name is be duplicated between database and mid-tier servers.

The "basic" and "asof" demonstrations should be reviwed before running these exercises. Serveral concepts introduced in those exercises are not explained here. Exercise #1 in this demonstration is similar to Exercise #1 in the asof demonstration. Additionally, the "DML & API Calls" and "Other Object Location" diagrams in the "DTGen_Notes.pdf" document in the "docs" directory provide a graphical layout of the multi-tier deployments created by DTGen.

This demonstration requires the creation of a database link, which uses Oracle's loopback to the same database. Configuration of this loopback is a normal part of the installation of an Oracle database. If this loopback is not working or has been disabled, these exercises will return errors.

The exercises in this directory are numbered and must be executed in sequential order. The demonstration users must be created with the "create_demo_users.sql" script in the parent directory before the first exercise is run. The demonstration users must be dropped with the "drop_demo_users.sql" script before the "create_demo_users.sql" script can be re-run. These exercises also assume that the default username/password (dtgen/dtgen) is still in use for the generator. Names and passwords are set in the "vars.sql" script and can be modified, if necessary. Also, the DTGen database objects must be installed in the database and the DTGen must be ready to generate code.

Exercise #1: Simple Mid-Tier

Command Line:

```
sqlplus /nolog @e1
```

Exercise #1 modifies the database. The "drop_demo_users.sql" and "create_demo_users.sql" scripts must be used to reset the database before re-running this exercise.

Based on the demobld.sql script, this exercise implements the EMP and DEPT tables using DTGen. The script for this exercise performs the following functions:

1. Removes any old DEMO1 Items from DTGEN
2. Creates new DEMO1 Items in DTGEN
3. Generates the DEMO1 Application in DTGEN
4. Creates the "install_db.sql" and "install_mt.sql" scripts
5. Runs the "install_db.sql" script

6. Loads Data and Shows EXPLAIN PLAN
7. Runs the "install_mt.sql" script
8. Shows EXPLAIN PLAN and Reports Data

Steps 1-3 are captured in the "e1.LST" file. Following is a example of e1.LST.

```
Login to dtgen
Connected.
Remove old DEMO3 Schema from DTGEN
create a DEMO3 Schema in DTGEN
Generate Demo3 Application
Capture SQL Scripts
```

Step 4 is captured in the "install_db.sql" and "install_mt.sql" files. These files have a combined size of about 216 kbytes and over 8,000 lines. Due to their size, they are not listed here. They contain all the code generated by DTGen for this exercise.

Steps 5 though 8 are captured in the "install.LST" file. Step 5 is the execution of the install_db.sql script.

```
Login to dtgen_db_demo
Connected.
```

```
FILE_NAME
-----
-) create_glob
```

```
FILE_NAME
-----
-) create_ods
```

```
TABLE_NAME
-----
*** dept ***
```

```
TABLE_NAME
-----
*** emp ***
```

```
FILE_NAME
-----
-) create_integ
```

```
TABLE_NAME
-----
*** dept ***
```

```
TABLE_NAME
-----
*** emp ***
```

```
FILE_NAME
-----
-) create_oltp
```

```
TABLE_NAME
-----
*** dept ***
```

```
TABLE_NAME
-----
*** emp ***
```

```
FILE_NAME
-----
-) create_aa
```

```
TABLE_NAME
-----
```

```

*** dept ***

TABLE_NAME
-----
*** emp ***

FILE_NAME
-----
-) create_mods

```

The above listing represents a successful installation of the simulated database tier in the "dtgen_db_demo" schema. The application is small in that it only has 2 tables. The data is loaded silently (Step 6). An explain plan for a simple query is run to demonstrate the query runs locally.

```

SQL>
SQL> explain plan set statement_id = 'D3_E1_Q1'
2      into plan_table for select * from emp where empno = 7900;
SQL>
SQL> select plan_table_output from table (
2      dbms_xplan.display('PLAN_TABLE', 'D3_E1_Q1') );

PLAN_TABLE_OUTPUT
-----
Plan hash value: 750405628

-----
| Id | Operation                                | Name | Rows | Bytes | Cost (%CPU) | Time      |
-----
|  0 | SELECT STATEMENT                        |      |     1 |    146 |      2   (0) | 00:00:01 |
|  1 |   TABLE ACCESS BY INDEX ROWID         | EMP  |     1 |    146 |      2   (0) | 00:00:01 |
|*  2 |    INDEX UNIQUE SCAN                   | EMP_NK |     1 |          |      1   (0) | 00:00:01 |
-----

Predicate Information (identified by operation id):
-----
  2 - access("EMPNO"=7900)

```

The above explain plan shows that the query of EMP has no REMOTE operations. This query was run locally. Step 7 is the execution of the install_mt.sql script below.

```

Login to dtgen_mt_demo
Connected.

FILE_NAME
-----
-) create_gdst

FILE_NAME
-----
-) create_dist

TABLE_NAME
-----
*** dept ***

TABLE_NAME
-----
*** emp ***

FILE_NAME
-----
-) create_oltp

TABLE_NAME
-----
*** dept ***

TABLE_NAME
-----
*** emp ***

FILE_NAME

```

```

-----
-) create_mods

FILE_NAME
-----
-) create_gdst_sec

FILE_NAME
-----
-) create_dist_sec

FILE_NAME
-----
-) create_oltp_sec

FILE_NAME
-----
-) create_mods_sec

```

The above listing represents a successful installation of the simulated mid-tier in the "dtgen_mt_demo". A database link is used to access the data in the tables at the simulated database tier ("dtgen_db_demo" schema). An explain plan for the same simple query below (step 8). It demonstrates the query is actually run remotely.

```

SQL> explain plan set statement_id = 'D3_E1_Q2'
2      into plan_table for select * from emp where empno = 7900;
SQL>
SQL> select plan_table_output from table (
2      dbms_xplan.display('PLAN_TABLE', 'D3_E1_Q2') );

```

PLAN_TABLE_OUTPUT

Plan hash value: 750405628

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Inst
0	SELECT STATEMENT REMOTE		1	146	2 (0)	00:00:01	
1	TABLE ACCESS BY INDEX ROWID	EMP	1	146	2 (0)	00:00:01	XE
* 2	INDEX UNIQUE SCAN	EMP_NK	1		1 (0)	00:00:01	XE

Predicate Information (identified by operation id):

2 - access("A1"."EMPNO"=7900)

Note

- fully remote statement

The explain plan above is identical to the previous explain plan, except for the note showing that it is a "fully remote statement". This confirms that the simulated tiers are working. The queries below also confirm that the data is accessible from the simulated mid-tier.

```

SQL>
SQL> select deptno, id, dname, loc, aud_beg_usr, aud_beg_dtm
2      from dept_act order by deptno, id;

```

DEPTNO	ID	DNAME	LOC	AUD_BEG_	AUD_BEG_D
10	1	ACCOUNTING	NEW YORK	Dataload	01-NOV-80
20	2	RESEARCH	DALLAS	Dataload	01-NOV-80
30	3	SALES	CHICAGO	THOMPSON	17-AUG-82
40	4	OPERATIONS	BOSTON	JAMES	12-FEB-82

```

SQL>
SQL> select deptno, dept_id, dname, loc, aud_beg_usr,
2      aud_beg_dtm, aud_end_usr, aud_end_dtm
3      from dept_aud
4      order by deptno, dept_id, aud_beg_dtm;

```

DEPTNO	DEPT_ID	DNAME	LOC	AUD_BEG_	AUD_BEG_D	AUD_END_	AUD_END_D
--------	---------	-------	-----	----------	-----------	----------	-----------

```

20      3 SALES      ST LOUIS Dataload 01-NOV-80 THOMPSON 17-AUG-82
40      4 OPERATIONS BUFFALO  Dataload 01-NOV-80 JAMES   12-FEB-82

```

SQL>

```

SQL> select empno, id, ename, job, mgr_emp_nkl, hiredate,
2      sal, dept_nkl, aud_beg_usr, aud_beg_dtm
3      from emp_act
4      order by empno, id;

```

EMPNO	ID	ENAME	JOB	MGR_	HIREDATE	SAL	DEPT_	AUD_BEG_	AUD_BEG_D
7369	3	SMITH	CLERK	7902	17-DEC-80	800	20	SMITH	26-FEB-83
7499	4	ALLEN	SALESMAN	7698	20-FEB-81	1600	30	THOMPSON	12-MAY-81
7521	5	WARD	SALESMAN	7698	22-FEB-81	1250	30	THOMPSON	14-MAY-81
7566	6	JONES	MANAGER	7839	02-APR-81	2975	20	SMITH	29-NOV-81
7654	14	MARTIN	SALESMAN	7698	28-SEP-81	1250	30	SMITH	26-SEP-81
7698	8	BLAKE	MANAGER	7839	01-MAY-81	2850	30	SMITH	30-NOV-81
7782	9	CLARK	MANAGER	7839	09-JUN-81	2450	10	SMITH	26-NOV-81
7788	20	SCOTT	ANALYST	7566	09-DEC-82	3000	20	JAMES	11-DEC-82
7839	15	KING	PRESIDENT		17-NOV-81	5000	10	SMITH	18-NOV-81
7844	13	TURNER	SALESMAN	7698	08-SEP-81	1500	30	SMITH	06-SEP-81
7876	21	ADAMS	CLERK	7788	12-JAN-83	1100	20	SMITH	12-JAN-83
7900	17	JAMES	CLERK	7698	03-DEC-81	950	30	SMITH	05-DEC-81
7902	18	FORD	ANALYST	7566	03-DEC-81	3000	20	SMITH	01-DEC-81
7934	19	MILLER	CLERK	7782	23-JAN-82	1300	10	JAMES	21-JAN-82

SQL>

```

SQL> select empno, emp_id, ename, job, mgr_emp_id,
2      aud_beg_usr, aud_beg_dtm, aud_end_usr, aud_end_dtm
3      from emp_hist
4      order by empno, emp_id, aud_beg_dtm;

```

EMPNO	EMP_ID	ENAME	JOB	MGR_	AUD_BEG_	AUD_BEG_D	AUD_END_	AUD_END_D
7301	1	ELLISON	PRESIDENT		DAVIS	04-NOV-80	THOMPSON	28-JUN-81
7344	2	DAVIS	CLERK	1	DAVIS	14-NOV-80	THOMPSON	25-JUN-81
7344	2	DAVIS	CLERK	11	THOMPSON	25-JUN-81	THOMPSON	20-AUG-81
7344	2	DAVIS	CLERK	12	THOMPSON	20-AUG-81	SMITH	29-NOV-81
7344	2	DAVIS	CLERK	15	SMITH	29-NOV-81	SMITH	06-DEC-81
7369	3	THOMPSON	CLERK	1	DAVIS	15-DEC-80	THOMPSON	25-JUN-81
7369	3	THOMPSON	CLERK	11	THOMPSON	25-JUN-81	SMITH	21-AUG-81
7369	3	SMITH	CLERK	12	SMITH	21-AUG-81	SMITH	01-DEC-81
7369	3	SMITH	CLERK	15	SMITH	01-DEC-81	SMITH	26-FEB-83
7499	4	ALLEN	SALESMAN	1	THOMPSON	17-FEB-81	THOMPSON	12-MAY-81
7521	5	WARD	SALESMAN	1	THOMPSON	24-FEB-81	THOMPSON	14-MAY-81
7566	6	JONES	MANAGER	1	THOMPSON	03-APR-81	THOMPSON	24-JUN-81
7566	6	JONES	MANAGER	11	THOMPSON	24-JUN-81	SMITH	22-AUG-81
7566	6	JONES	MANAGER	12	SMITH	22-AUG-81	SMITH	29-NOV-81
7654	7	MARTIN	SALESMAN	1	THOMPSON	16-APR-81	THOMPSON	13-MAY-81
7698	8	BLAKE	MANAGER	1	THOMPSON	02-MAY-81	THOMPSON	24-JUN-81
7698	8	BLAKE	MANAGER	11	THOMPSON	24-JUN-81	THOMPSON	19-AUG-81
7698	8	BLAKE	MANAGER	12	THOMPSON	19-AUG-81	SMITH	30-NOV-81
7782	9	CLARK	MANAGER	1	THOMPSON	07-JUN-81	THOMPSON	23-JUN-81
7782	9	CLARK	MANAGER	11	THOMPSON	23-JUN-81	SMITH	23-AUG-81
7782	9	CLARK	MANAGER	12	SMITH	23-AUG-81	SMITH	26-NOV-81
7788	10	SCOTT	ANALYST	6	THOMPSON	10-JUN-81	JAMES	09-MAR-82
7839	11	KING	PRESIDENT		THOMPSON	18-JUN-81	SMITH	30-AUG-81
7840	12	LANE	PRESIDENT		THOMPSON	12-AUG-81	SMITH	29-NOV-81
7876	16	ADAMS	CLERK	6	SMITH	20-NOV-81	JAMES	13-JUN-82

With the completion of exercise 1, a new application was defined in DTGen, generated, and loaded into a single database using 2 different schema names, which correctly simulates a simple mid-tier.

Exercise #2: Materialized Views

Command Line:

```
sqlplus /nolog @e2
```

Exercise #2 does not modify the database. This exercise can be repeated without problem.

In the exercise #2, a basic generation and installation of a simple mid-tier was completed. Part of that installation was a materialized view for the DEPT table on the simulated mid-tier.

```
Login to dtgen
Connected.
```

```
SQL> select seq, name, type, mv_refresh_hr from tables_act
2   where applications_nk1 = 'DEMO3' order by seq;
```

SEQ	NAME	TYP	MV_REFRESH_HR
10	dept	LOG	24
20	emp	EFF	

The results above show the DTGen configuration to setup a materialized view. The value in the MV_REFRESH_HR field signals the generator to make this table a materialized view and identifies the refresh period in hours. The materialized view on the simulated mid-tier will be refreshed daily.

```
Login to dtgen_mt_demo
Connected.
```

```
SQL>
SQL> select mview_name, last_refresh_date from user_mviews;
```

MVIEW_NAME	LAST_REFR
DEPT	23-APR-12

The results above confirm that a materialized view exists in the simulated mid-tier schema. Below is an explain plan on a query of that view.

```
SQL>
SQL> explain plan set statement_id = 'D3_E2_Q1'
2   into plan_table for select * from dept where deptno = 40;
```

```
SQL>
SQL> select plan_table_output from table (
2   dbms_xplan.display('PLAN_TABLE', 'D3_E2_Q1') );
```

```
PLAN_TABLE_OUTPUT
```

```
Plan hash value: 1799509361
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	73	1 (0)	00:00:01
1	MAT_VIEW ACCESS BY INDEX ROWID	DEPT	1	73	1 (0)	00:00:01
* 2	INDEX UNIQUE SCAN	DEPT_NK	1		0 (0)	00:00:01

```
Predicate Information (identified by operation id):
```

```
2 - access("DEPTNO">=40)
```

Step 2 in the PLAN_TABLE_OUTPUT above show a MAT_VIEW, or materialized view, access operation was performed. The data returned for this query is local to the simulated mid-tier, contrasted with the EMP query in exercise #1 that returned data from the remote simulated database tier. Also note that DTGen automatically generated the index DEPT_NK on the materialized view.

Exercise #3: User Security

Command Line:

```
sqlplus /nolog @e3
```

Exercise #3 modifies the database. The "drop_demo_users.sql", "create_demo_users.sql", and "e1.sql" scripts must be used to reset the database before re-running this exercise.

In this exercise, the user security layer is demonstrated. The security layer is implemented by separating the user login from the schema that contains the database objects. Synonyms are used to easily access the database objects in the other schema. Below is a query showing the DTGen configuration that identifies the name of the schema with that contains the database objects

```
Login to dtgen
Connected.
```

```
SQL> select name, db_schema from applications where abbr = 'DEMO3';
```

NAME	DB_SCHEMA
Dtgen Tiers Demonstration	DTGEN_MT_DEMO

From the configuration listed above an "install_usr.sql" script was created in Exercise #1 that defines the required synonyms. This exercise continues with the login and installation of the user synonyms in the install_usr.sql script.

```
Login to dtgen_usr_demo
Connected.
```

```
FILE_NAME
-----
-) create_usyn
```

```
TABLE_NAME
-----
*** dept ***
```

```
TABLE_NAME
-----
*** emp ***
```

This listing above is a successful installation of the user synonyms. An explain plan for the same simple query run in Exercise #1 demonstrates the query is actually run remotely.

```
SQL> explain plan set statement_id = 'D3_E3_Q1'
2 into plan_table for select * from emp where empno = 7900;
```

```
Explained.
```

```
SQL> select plan_table_output from table (
2 dbms_xplan.display('PLAN_TABLE', 'D3_E3_Q1') );
```

```
PLAN TABLE OUTPUT
-----
Plan hash value: 750405628
-----
|Id|Operation                               |Name|Rows|Bytes|Cost(%CPU)|Time|Inst|
-----|-----|-----|-----|-----|-----|-----|
| 0|SELECT STATEMENT REMOTE                  |    |1|146|2 (0)|00:00:01|
| 1|TABLE ACCESS BY INDEX ROWID EMP          |EMP |1|146|2 (0)|00:00:01|XE|
|* 2|INDEX UNIQUE SCAN                       |EMP_NK|1| |1 (0)|00:00:01|XE|
-----
```


Predicate Information (identified by operation id):

2 - access("A1"."EMPNO"=7900)

Note

- fully remote statement

The explain plan above is identical to the previous explain plan, except for the note showing that it is a "fully remote statement". This confirms that the simulated tiers are working. A query of the ASOF data is run to demonstrate user access to the database objects in the other schema.

```
SQL>
SQL> execute glob.set_asof_dtm(to_timestamp('1981-06-01', 'YYYY-MM-DD'))
SQL>
SQL> select empno, ename, job, mgr_emp_nk1, hiredate, sal, deptno, dname, loc
       2   from emp_asof e, dept_asof d where e.dept_id = d.id
       3   order by empno;
```

EMPNO	ENAME	JOB	MGR_EMP_NK1	HIREDATE	SAL	DEPTNO	DNAME	LOC
7301	ELLISON	PRESIDENT		02-NOV-80	4000	10	ACCOUNTING	NEW YORK
7344	DAVIS	CLERK	7301	16-NOV-80	1400	10	ACCOUNTING	NEW YORK
7369	THOMPSON	CLERK	7301	17-DEC-80	800	10	ACCOUNTING	NEW YORK
7499	ALLEN	SALESMAN	7698	20-FEB-81	1600	20	SALES	ST LOUIS
7521	WARD	SALESMAN	7698	22-FEB-81	1250	20	SALES	ST LOUIS
7566	JONES	MANAGER	7301	02-APR-81	2975	20	RESEARCH	DALLAS
7698	BLAKE	MANAGER	7301	01-MAY-81	2850	20	SALES	ST LOUIS

With the execution of the above query, data was successfully retrieved from a remote database link (database/mid-tier server tiers) via a synonym pointed to another schema (database user/schema tiers).

Exercise #4: Global Locks

Command Line:

```
sqlplus /nolog @e4
```

Exercise #5 does not modify the database. This exercise can be repeated without problem.

In an application that spans multiple nodes, it is sometimes necessary to single-thread a process or ensure sole access to an global resource. This functionality is performed by the GLOB.package.

```
Login to dtgen_usr_demo
Connected.
```

```
SQL>
SQL> describe glob
```

FUNCTION GET_ASOF_DTM RETURNS TIMESTAMP WITH TIME ZONE

FUNCTION GET_DB_CONSTRAINTS RETURNS BOOLEAN

FUNCTION GET_DTM RETURNS TIMESTAMP WITH LOCAL TIME ZONE

FUNCTION GET_FOLD_STRINGS RETURNS BOOLEAN

FUNCTION RELEASE_LOCK RETURNS VARCHAR2

FUNCTION REQUEST_LOCK RETURNS VARCHAR2

Argument Name	Type	In/Out	Default?
LOCKNAME_IN	VARCHAR2	IN	

Argument Name	Type	In/Out	Default?
ASOF_DTM_IN	TIMESTAMP WITH TIME ZONE	IN	

Argument Name	Type	In/Out	Default?
BOOL_IN	BOOLEAN	IN	

Argument Name	Type	In/Out	Default?
BOOL_IN	BOOLEAN	IN	

Argument Name	Type	In/Out	Default?
TABLE_NAME	VARCHAR2	IN	
EFF_DTM_IN	TIMESTAMP	IN	

The output above from e4.LST has a describe command that shows the names and parameters of all the available functions and procedures in the GLOB package. The next portion of the e4.sql script will request a lock.

```
SQL>
SQL> execute util.set_usr('DEMO3');
SQL> execute dbms_output.put_line(glob.request_lock('DEMO3'));
SUCCESS
SQL>
SQL> set echo off
```

Open another window to run "e4a.sql"
Ex: sqlplus /nolog @e4a

Press Enter when prompted from other window

In the output above, the GLOB.REQUEST_LOCK function returned the status of SUCCESS. The session at the database tier (simulated) is now holding the lock for the session at the mid-tier (simulated). A pause in the e4.sql script allows another mid-tier (and database tier) session to be started. The same lock is requested from the other session using the e4a.sql script.

```
Login to dtgen_usr_demo
Connected.

SQL> execute util.set_usr('DEMO3');
SQL> execute dbms_output.put_line(glob.request_lock('DEMO3', 2));
TIMEOUT

SQL> -- Press Enter on the other window to continue
SQL> execute dbms_output.put_line(glob.request_lock('DEMO3'));
```

In the output above from e4a.LST, the first lock request failed with a timeout of 2 seconds. The second request halts the execution of the script, waiting to acquire the lock without a timeout value set. Moving back to the initial window, ENTER is pressed to continue the execution of the e4.sql script.

```
SQL> execute dbms_output.put_line(glob.request_lock('DEMO2'));
RELEASE ONLY

SQL> execute dbms_output.put_line(glob.request_lock('DEMO3'));
SUCCESS

SQL> execute dbms_output.put_line(glob.release_lock);
SUCCESS

SQL> execute dbms_output.put_line(glob.release_lock);
SUCCESS
```

A request for a lock, other than the one already held, is executed and fails. The RELEASE ONLY return confirms that the first lock must be released before another lock can be aquired. A request for the the lock that is already held succeeds. Both lock release calls succeed, the first because the lock is actually released, and the second because not lock is currently held by the session. As soon as the lock was release, execution continued on e4a.sql in the other window.

```
SUCCESS
```

```
SQL> execute dbms_output.put_line(glob.release_lock);  
SUCCESS
```

The lock that was held by the initial window is successfully requested and released. This locking mechanism works amoung all processes running on any database tier node (RAC) or any mid-tier node.